

Ana Stanojevic

Zürich, Switzerland • contact@ana-stanojevic.com • ana-stanojevic.com • github.com/ana-stanojevic

Applied AI Engineer — LLM Systems, Agents & AI Automation

Applied AI Engineer building production-oriented LLM systems, agentic workflows, retrieval architectures, and AI automation pipelines across Huawei, IBM, Google, and EPFL.

EXPERIENCE

Sabbatical - Independent AI Engineering & Systems Development

03/2025 - Present

Bounded Job Application Agent

- Autonomous AI system for job discovery, ranking, preparation, and browser-assisted application execution.
- Designed and built production-oriented agentic AI workflows integrating retrieval, orchestration, structured generation, tool execution, evaluation, and browser automation.
- Developed scalable AI infrastructure using FastAPI, async Python, Docker, vector retrieval, modular APIs, and distributed workflow patterns.
- Implemented semantic retrieval, confidence-based routing, validation, retries, observability, and execution pipelines for autonomous AI workflow automation.

Huawei

01/2024 - 03/2025

AI / LLM Systems Engineer

Zürich, Switzerland

- Designed and optimized production-oriented RAG pipelines using FAISS-based semantic retrieval.
- Built scalable retrieval and inference systems focused on low-latency semantic search and runtime reliability.
- Developed modular AI components for orchestration, evaluation, and execution workflows.
- Worked on distributed LLM inference systems and transformer runtime optimization.
- Contributed to scalable AI serving infrastructure and production-oriented LLM workflow execution under real-world deployment constraints.

IBM

10/2019 - 12/2023

Machine Learning Engineer

Zürich, Switzerland

- Developed scalable ML systems and inference methods for spiking neural networks under performance and deployment constraints.
- Built production-oriented computer vision prototypes and optimized neural inference execution.
- Worked on efficient neural architectures, runtime optimization, and scalable ML workflows.
- Publications in Nature Communications and Neural Networks, patents in neural inference optimization.

Google

08/2018 - 02/2019

Software Engineer Intern (Google Assistant)

Zürich, Switzerland

- Worked with BERT-based NLP systems in C++ and Python.
- Contributed to conversational NLP workflows and answer relevance optimization for large-scale assistant systems.

EDUCATION

PhD in Computer & Communication Sciences - EPFL

MSc in Communication Systems (Data Analytics) - EPFL

BSc in Electrical & Computer Engineering - University of Belgrade

SKILLS

AI Systems & GenAI: AI agents, agentic workflows, browser automation, embeddings, evaluation pipelines, LLM systems, OpenAI API, Anthropic API, orchestration frameworks, RAG architectures, semantic search, structured generation, structured outputs, tool calling, MCP, vector databases (FAISS)

Engineering & Infrastructure: async Python, CI/CD, distributed inference, Docker, FastAPI, Git, Linux, Playwright, modular backend systems, observability, Python, REST APIs, SQL, structured logging

ML & Frameworks: Apache Spark, CUDA, Hugging Face, Kafka, LangChain, PyTorch, scikit-learn, TensorFlow, Transformers

Languages: English (C2), German (B2), Serbian (native)

AWARDS

Pat Goldberg Honorable Mention Award (IBM Research Zürich) • Excellence Fellowship (EPFL) • Dositeja Award • ETH Zürich Summer School on Open Source IC Design